

QUANTITATIVE ANALYSIS

VERBS FOR CLEANING DATA

AGENDA

1. Tidy Data: A Review
2. dplyr Verbs
3. Piping Functions

1 TIDY DATA: A REVIEW

KEY CHARACTERISTICS

F	M	A
↑	↑	↑
↓	↓	↓

Each variable is stored
in its own column

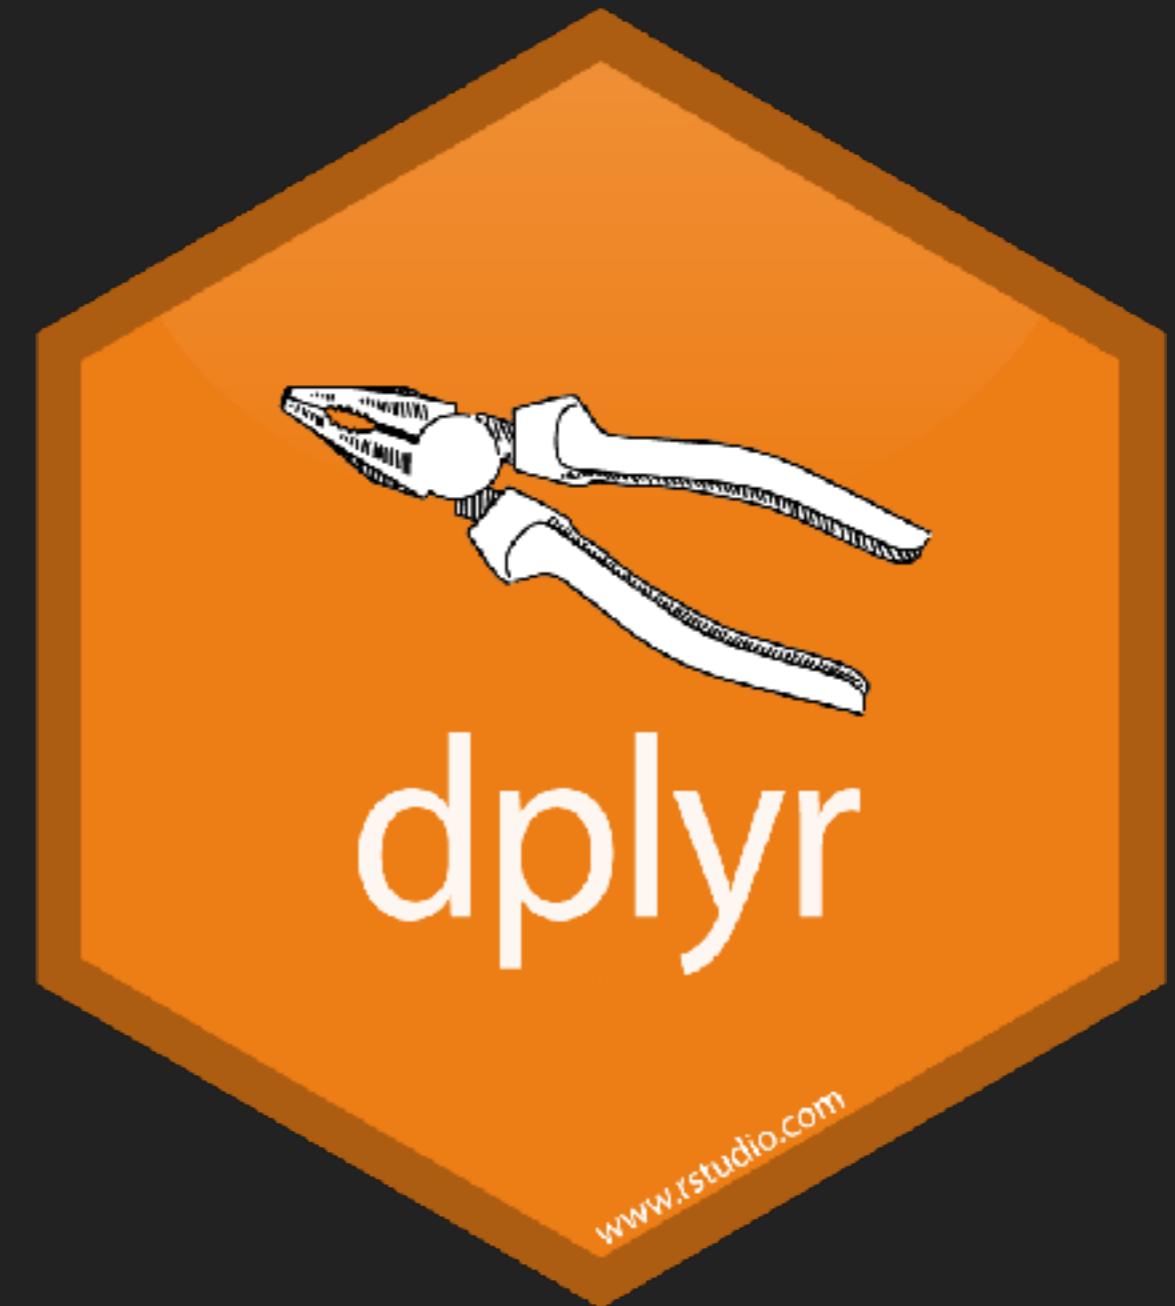
F	M	A
←→		
←→		
←→		

Each observation is stored
in its own row

2 DPLYR VERBS

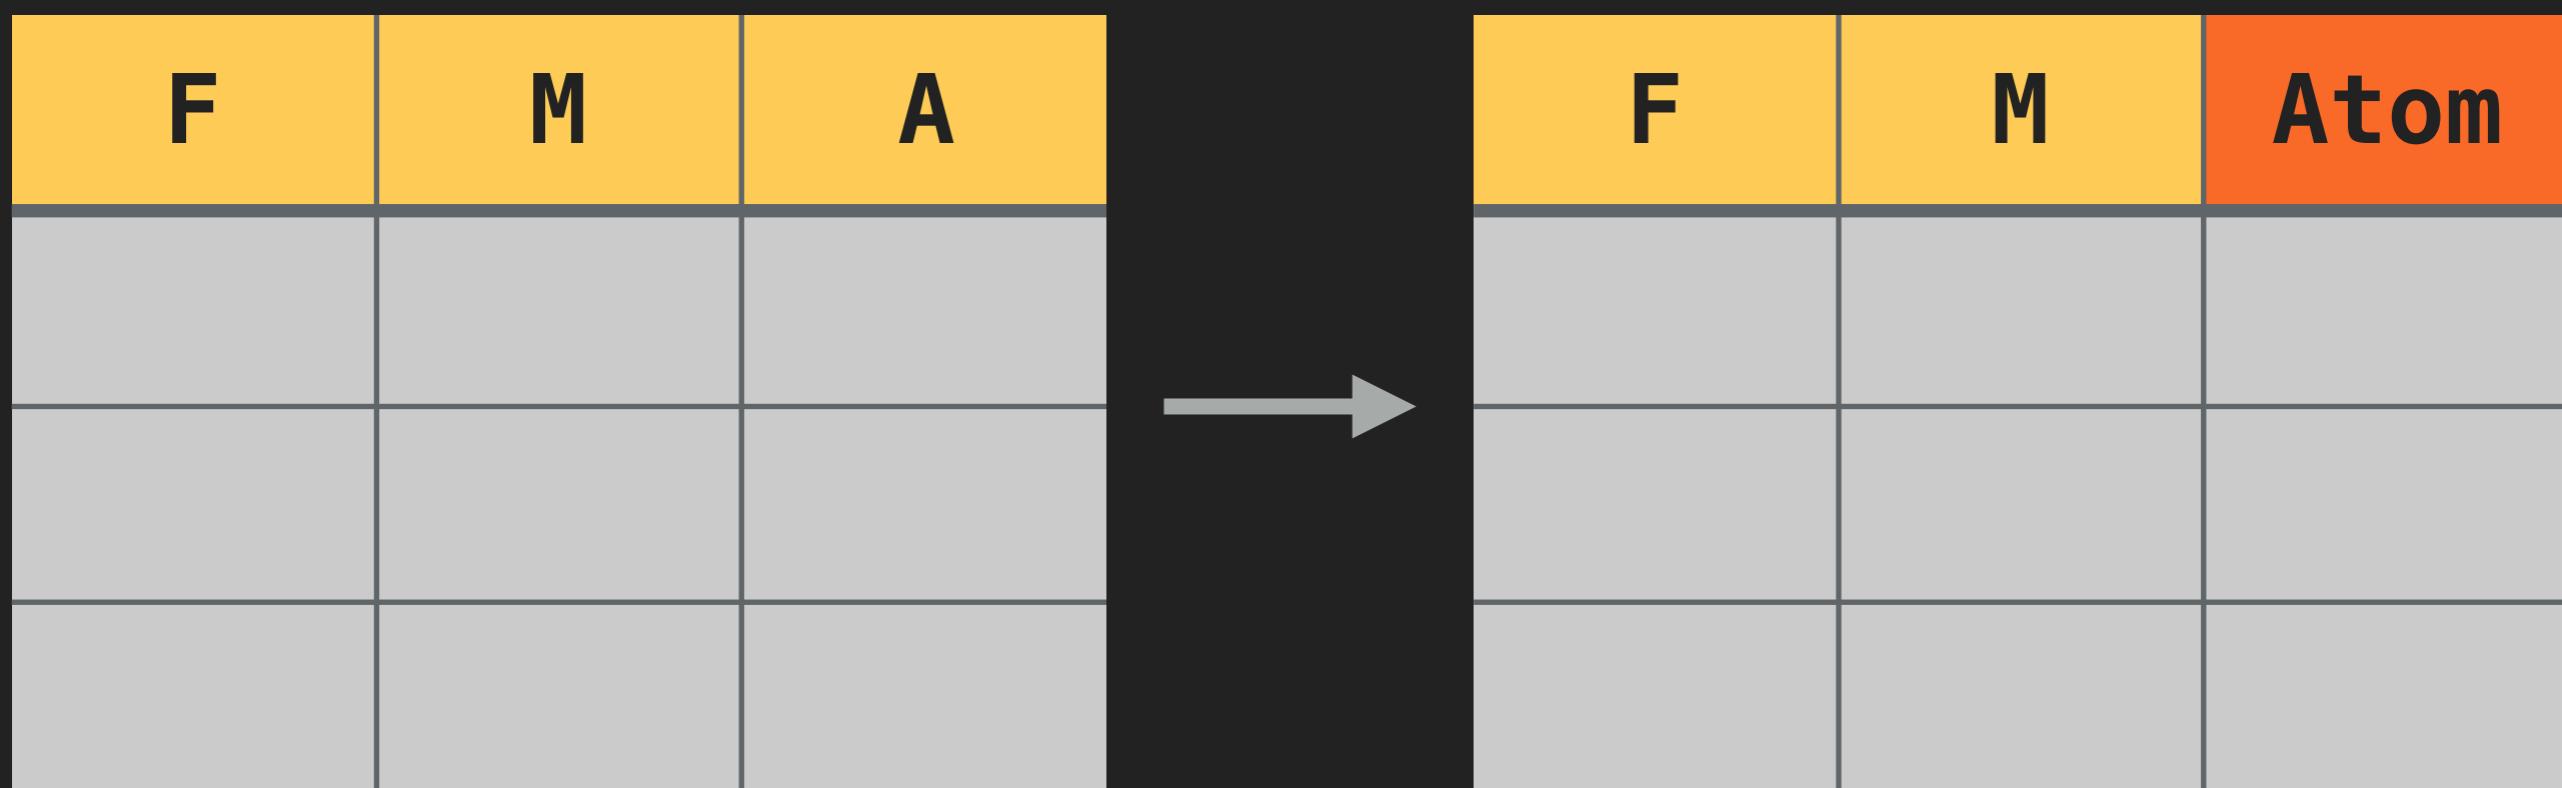
DPLYR

- ▶ Like `ggplot2`, `dplyr` is a core part of the `tidyverse`.
- ▶ `Dplyr` specializes in data wrangling, which is the work we put into getting a data set ready for analysis
- ▶ It is based around the concept of *verbs* - functions are named for actions that they undertake
- ▶ We'll focus on five key functions today



2. DPLYR VERBS

RENAMEING VARIABLES



RENAMING VARIABLES

```
dplyr::rename(dataFrame, newName = oldName)
```



Example - the mpg data from ggplot2:

```
rename(mpg, hwyMpg = hwy)
```

RENAMING VARIABLES

```
dplyr::rename(dataFrame, newName = oldName)
```



Example - the mpg data from ggplot2:

```
rename(mpg, hwyMpg = hwy)
```



This does not make the change permanent, however. You must *assign* the results of dplyr functions back to the original data frame or to a new one.

ASSIGNING CHANGES

```
dataFrame <- rename(dataFrame, newName = oldName)
```

-  Example 1 - assigning the mpg data from ggplot2 to a new object:

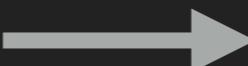
```
autoData <- rename(mpg, hwyMpg = hwy)
```

-  Example 2 - overwriting the autoData data example 1:

```
autoData <- rename(autoData, type = class)
```

REORDERING OBSERVATIONS

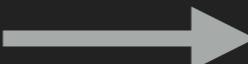
F	M	A
2		
5		
3		
4		
8		
1		



F	M	A
1		
2		
3		
4		
5		
8		

REORDERING OBSERVATIONS

F	M	A
2		
5		
3		
4		
8		
1		



F	M	A
8		
5		
4		
3		
2		
1		

REORDERING OBSERVATIONS

`dplyr::arrange(dataFrame, varlist)`

-  Example - the mpg data from ggplot2 in *ascending* order (lowest first):

`arrange(mpg, hwy)`

-  You can include more than one variable, separated by commas, if you want a list sorted based on more than one condition.

Reordering your data may change how some output looks and how the assignment of ID numbers occurs.

REORDERING OBSERVATIONS

```
dplyr::arrange(dataFrame, desc(varlist))
```



Example - the mpg data from ggplot2 in *descending* order (highest first):

```
arrange(mpg, desc(hwy))
```



You can include more than one variable, separated by commas, if you want a list sorted based on more than one condition.

Reordering your data may change how some output looks and how the assignment of ID numbers occurs.

REORDERING OBSERVATIONS

```
> library(tidyverse)

> autoData <- mpg

> head(autoData)
# A tibble: 6 × 11
  manufacturer model displ year cyl trans drv cty hwy fl class
  <chr> <chr> <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
1 audi     a4      1.8  1999     4 auto(l5) f    18    29 p   compact
2 audi     a4      1.8  1999     4 manual(m5) f    21    29 p   compact
3 audi     a4      2.0  2008     4 manual(m6) f    20    31 p   compact
4 audi     a4      2.0  2008     4 auto(av)   f    21    30 p   compact
5 audi     a4      2.8  1999     6 auto(l5)   f    16    26 p   compact
6 audi     a4      2.8  1999     6 manual(m5) f    18    26 p   compact

> View(autoData)
```

REORDERING OBSERVATIONS

```
> tail(autoData)
# A tibble: 6 x 11
  manufacturer model displ year cyl trans drv cty hwy fl class
  <chr>     <chr> <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>
1 volkswagen passat    1.8  1999     4 auto(l5)   f    18    29   p midsize
2 volkswagen passat    2.0  2008     4 auto(s6)   f    19    28   p midsize
3 volkswagen passat    2.0  2008     4 manual(m6) f    21    29   p midsize
4 volkswagen passat    2.8  1999     6 auto(l5)   f    16    26   p midsize
5 volkswagen passat    2.8  1999     6 manual(m5) f    18    26   p midsize
6 volkswagen passat    3.6  2008     6 auto(s6)   f    17    26   p midsize
```

REORDERING OBSERVATIONS

```
> autoData <- arrange(autoData, hwy)

> head(autoData)
# A tibble: 6 x 11
  manufacturer      model  displ   year   cyl     trans   drv     cty     hwy     fl class
  <chr>            <chr>  <dbl> <int> <int>    <chr>   <chr> <int> <int> <chr> <chr>
1 dodge           dakota pickup  4.7  2008     8 auto(l5)  4       9     12   e  pickup
2 dodge           durango 4wd    4.7  2008     8 auto(l5)  4       9     12   e   suv
3 dodge          ram 1500 pickup  4.7  2008     8 auto(l5)  4       9     12   e  pickup
4 dodge          ram 1500 pickup  4.7  2008     8 manual(m6) 4       9     12   e  pickup
5 jeep            grand cherokee 4wd  4.7  2008     8 auto(l5)  4       9     12   e   suv
6 chevrolet       k1500 tahoe 4wd  5.3  2008     8 auto(l4)  4      11     14   e   suv
```

REORDERING OBSERVATIONS

```
> autoData <- arrange(autoData, desc(hwy))
```

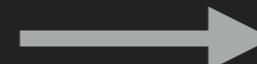
```
> head(autoData)
```

```
# A tibble: 6 x 11
```

	manufacturer	model	displ	year	cyl	trans	drv	cty	hwy	fl	class
	<chr>	<chr>	<dbl>	<int>	<int>	<chr>	<chr>	<int>	<int>	<chr>	<chr>
1	volkswagen	jetta	1.9	1999	4	manual(m5)	f	33	44	d	compact
2	volkswagen	new beetle	1.9	1999	4	manual(m5)	f	35	44	d	subcompact
3	volkswagen	new beetle	1.9	1999	4	auto(l4)	f	29	41	d	subcompact
4	toyota	corolla	1.8	2008	4	manual(m5)	f	28	37	r	compact
5	honda	civic	1.8	2008	4	auto(l5)	f	25	36	r	subcompact
6	honda	civic	1.8	2008	4	auto(l5)	f	24	36	c	subcompact

SUBSETTING DATA

F	M	A
	a	
	a	
b		b
b		b
b		b
	a	



F	M	A
	b	
	b	
b		b

SUBSETTING DATA

`dplyr::filter(dataFrame, expression)`

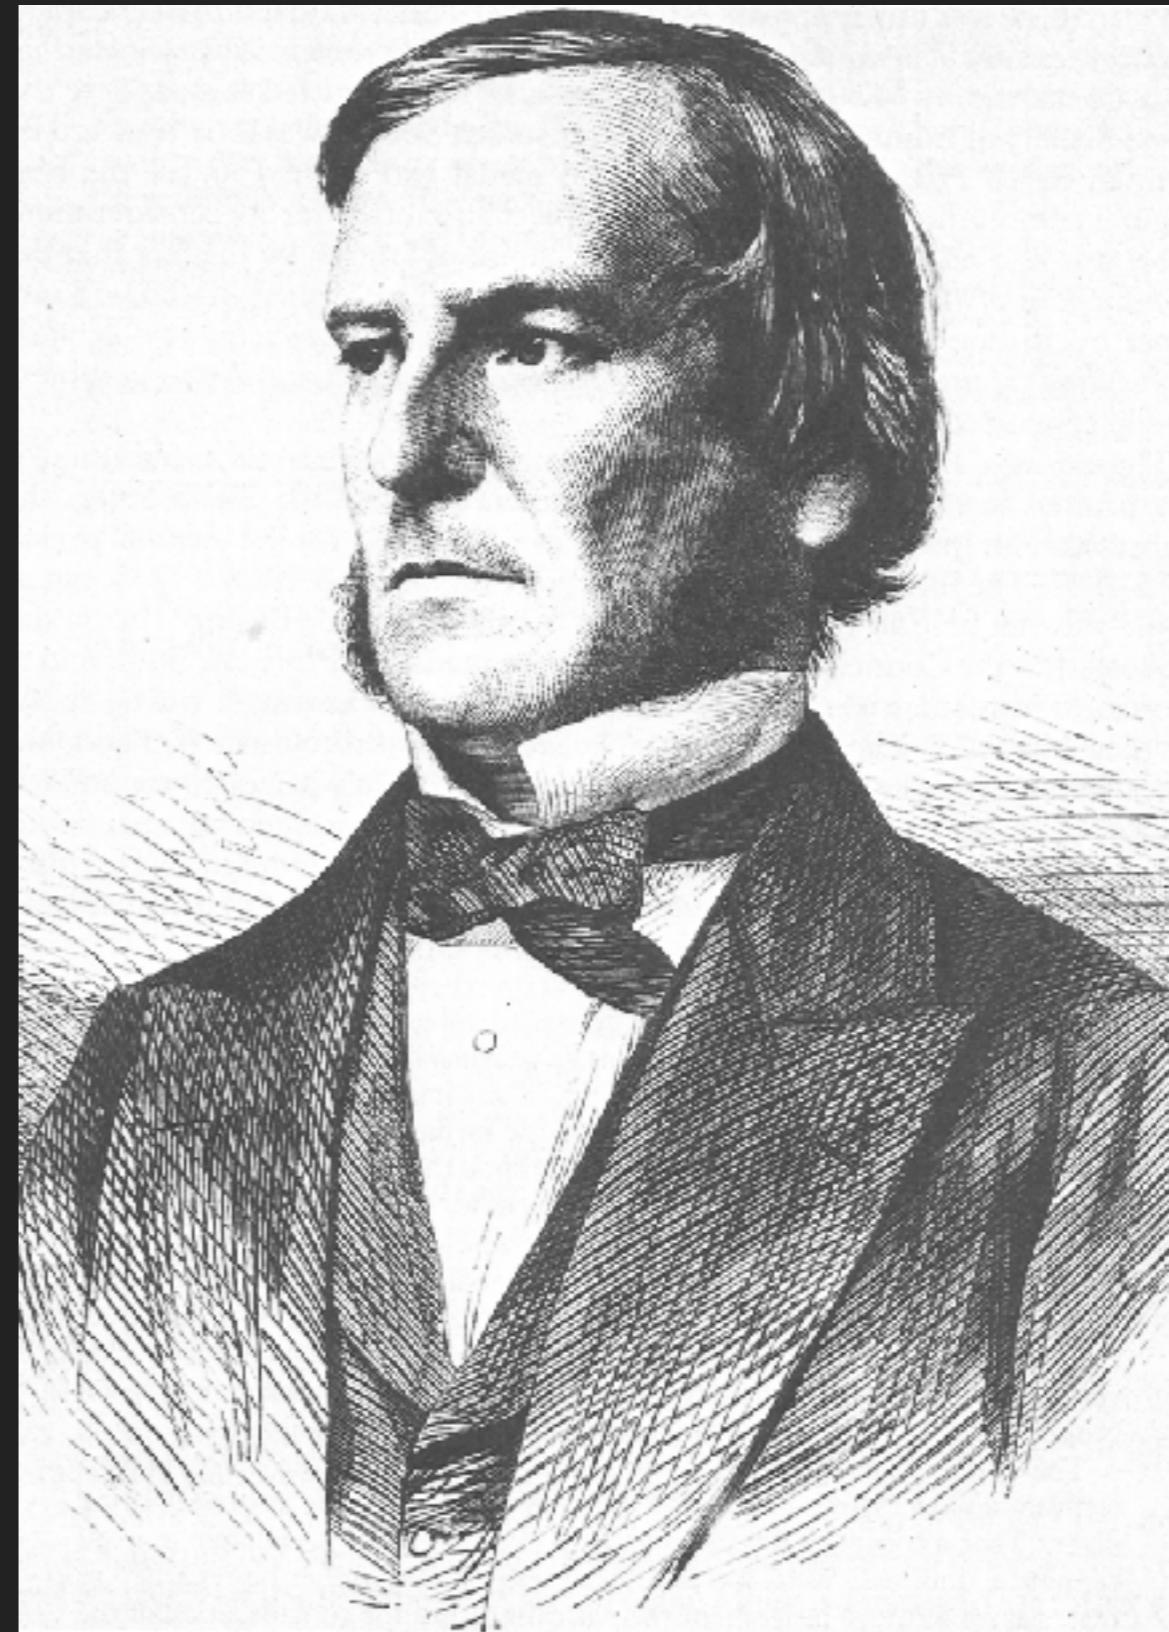
-  Example - the mpg data from ggplot2 filtered using a numeric value:

```
filter(mpg, hwy >= 30)
```

-  This will *retain* only observations that are TRUE based on the expression.

GEORGE BOOLE

- ▶ British mathematician who was active during the 1840s and 1850s
- ▶ Credited with establishing the field of boolean algebra in papers published in 1847 and 1854
- ▶ Boolean algebra is premised on the idea that logical relations can be used evaluate expressions as either TRUE or FALSE
- ▶ Boolean logic is a fundamental concept for modern computing



2. DPLYR VERBS

BOOLEAN LOGIC

```
filter(mpg, hwy >= 30)
```



model	year	hwy	boolean eval.
a4	1999	29	FALSE
forester awd	2008	23	FALSE
corolla	2008	35	TRUE



model	year	hwy
corolla	2008	35

SUBSETTING DATA

```
dplyr::filter(dataFrame, expression)
```



Example - the mpg data from ggplot2 filtered using a string:

```
filter(mpg, manufacturer == "subaru")
```



This will *retain* only observations that are TRUE based on the expression.

This method of searching strings is case sensitive and will only evaluate as TRUE for exact matches. There are more flexible ways to search strings as well.

SUBSETTING DATA

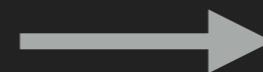
```
> library(tidyverse)

> subaru <- filter(mpg, manufacturer == "subaru")

> str(subaru)
Classes 'tbl_df', 'tbl' and 'data.frame': 14 obs. of 11 variables:
 $ manufacturer: chr  "subaru" "subaru" "subaru" "subaru" ...
 $ model        : chr  "forester awd" "impreza awd" "impreza awd" "forester awd" ...
 $ displ         : num  2.5 2.5 2.5 2.5 2.2 2.2 2.5 2.5 2.5 ...
 $ year          : int  2008 2008 2008 2008 1999 1999 1999 1999 2008 ...
 $ cyl           : int  4 4 4 4 4 4 4 4 4 ...
 $ trans         : chr  "manual(m5)" "auto(s4)" "manual(m5)" "auto(l4)" ...
 $ drv            : chr  "4" "4" "4" "4" ...
 $ cty            : int  20 20 20 20 21 19 19 19 18 19 ...
 $ hwy            : int  27 27 27 26 26 26 26 26 25 25 ...
 $ fl             : chr  "r" "r" "r" "r" ...
 $ class          : chr  "suv" "compact" "compact" "suv" ...
```

2. DPLYR VERBS

SUBSETTING DATA



A 6x6 grid of orange squares. The grid is bounded by a thick black vertical line on the left and a thick black horizontal line at the top. The grid is divided into six columns and six rows of equal size.

SUBSETTING DATA

```
dplyr::select(dataFrame, varlist)
```



Example - the mpg data from ggplot2:

```
select(mpg, manufacturer, model, hwy, class)
```



This approach will *retain* only the listed variables.

There are additional helper functions for searching

SUBSETTING DATA

```
> library(tidyverse)

> autoData <- select(mpg, manufacturer, model, hwy, class)

> str(autoData)
Classes 'tbl_df', 'tbl' and 'data.frame': 234 obs. of 4 variables:
 $ manufacturer: chr  "audi" "audi" "audi" "audi" ...
 $ model       : chr  "a4"   "a4"   "a4"   "a4"   ...
 $ hwy         : int  29  29  31  30  26  26  27  26  25  28 ...
 $ class        : chr  "compact" "compact" "compact" "compact" ...
```

SUBSETTING DATA

```
dplyr::select(dataFrame, -varlist)
```



Example - the mpg data from ggplot2:

```
select(mpg, -manufacturer, -model, -hwy, -class)
```



This approach will *remove* only the listed variables.

SUBSETTING DATA

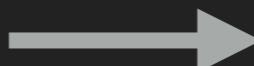
```
> library(tidyverse)

> autoData <- select(mpg, -manufacturer, -model, -hwy, -class)

> str(autoData)
Classes 'tbl_df', 'tbl' and 'data.frame': 234 obs. of 7 variables:
 $ displ: num 1.8 1.8 2 2 2.8 2.8 3.1 1.8 1.8 2 ...
 $ year : int 1999 1999 2008 2008 1999 1999 2008 1999 1999 2008 ...
 $ cyl  : int 4 4 4 4 6 6 6 4 4 4 ...
 $ trans: chr "auto(l5)" "manual(m5)" "manual(m6)" "auto(av)" ...
 $ drv   : chr "f" "f" "f" "f" ...
 $ cty   : int 18 21 20 21 16 18 18 18 16 20 ...
 $ fl    : chr "p" "p" "p" "p" ...
```

2. DPLYR VERBS

CREATING NEW VARIABLES



CREATING NEW VARIABLES

`dplyr::mutate(dataFrame, newVar = expression)`

-  Example - numerical calculation with the mpg data from ggplot2:

```
mutate(mpg, avgMpg = (cty+hwy)/2)
```

-  Requires numeric data

CREATING NEW VARIABLES

```
dplyr::mutate(dataFrame, newVar =  
  ifelse(expression, trueOutcome, falseOutcome))
```

 Example - binary variable creation with the mpg data from ggplot2:

```
mutate(mpg, highMpg = ifelse(hwy >= 30, TRUE, FALSE))
```



Requires numeric data.

True and false expressions can be either logical, character, or numeric data. You should be consistent in keeping both the true and false expressions as the same data type.

CREATING NEW VARIABLES

```
dplyr::mutate(dataFrame, newVar =  
  ifelse(expression, trueOutcome, falseOutcome))
```

 Example - binary variable creation with the mpg data from ggplot2:

```
mutate(mpg, subaru =  
  ifelse(manufacturer == "subaru", TRUE, FALSE))
```



Requires string data.

This method of searching strings is case sensitive and will only evaluate as TRUE for exact matches. There are more flexible ways to search strings as well.

CREATING NEW VARIABLES

```
> library(tidyverse)  
  
> autoData <- mpg  
  
> mutate(autoData, subaru = ifelse(manufacturer == "subaru", TRUE, FALSE))  
  
> table(autoData$subaru)  
  
FALSE   TRUE  
220     14
```

3 PIPING DATA

3. PIPING DATA

ASSIGNING DATA CAN GET CUMBERSOME

```
> library(tidyverse)

> japaneseAutos <- mpg

> japaneseAutos <-
  select(japaneseAutos, model, cty, hwy)

> japaneseAutos <-
  rename(japaneseAutos, cityMpg = cty)

> japaneseAutos <-
  rename(japaneseAutos, hwyMpg = hwy)

> japaneseAutos <-
  filter(japaneseAutos, manufacturer == "honda" |
  manufacturer == "nissan" |
  manufacturer == "subaru" |
  manufacturer == "toyota")

> japaneseAutos <-
  mutate(japaneseAutos, avgMpg =
  (cityMpg+hwyMpg)/2)

> japaneseAutos <- arrange(japaneseAutos, avgMpg)
```



LET US CHANGE OUR
TRADITIONAL ATTITUDE TO THE
CONSTRUCTION OF PROGRAMS:
INSTEAD OF IMAGINING THAT OUR
MAIN TASK IS TO INSTRUCT A
COMPUTER WHAT TO DO, LET US
CONCENTRATE RATHER ON
EXPLAINING TO HUMANS WHAT
WE WANT THE COMPUTER TO DO.

Donald E. Knuth

Stanford University Computer Scientist

3. PIPING DATA

MAGRITTR PACKAGE

- ▶ dplyr automatically loads the magrittr package
- ▶ magrittr includes a number of helpful functions, but is most well known for the “pipe”:

`%>%`



- ▶ Piping data makes it easier to write and more readable for humans

3. PIPING DATA

ASSIGNING DATA CAN GET CUMBERSOME

```
> library(tidyverse)  
  
> japaneseAutos <- mpg  
  
> japaneseAutos <-  
  select(japaneseAutos, model, cty, hwy)  
  
> japaneseAutos <-  
  rename(japaneseAutos, cityMpg = cty)  
  
> japaneseAutos <-  
  rename(japaneseAutos, hwyMpg = hwy)  
  
> japaneseAutos <-  
  filter(japaneseAutos, manufacturer == "honda" |  
  manufacturer == "nissan" |  
  manufacturer == "subaru" |  
  manufacturer == "toyota")  
  
> japaneseAutos <-  
  mutate(japaneseAutos, avgMpg =  
  (cityMpg+hwyMpg)/2)  
  
> japaneseAutos <- arrange(japaneseAutos, avgMpg)
```

```
> library(tidyverse)  
  
> mpg %>%  
  select(manufacturer,  
         model, cty, hwy) %>%  
  rename(cityMpg = cty) %>%  
  rename(hwyMpg = hwy) %>%  
  filter(manufacturer == "honda" |  
         manufacturer == "nissan" |  
         manufacturer == "subaru" |  
         manufacturer == "toyota") %>%  
  mutate(avgMpg =  
        (cityMpg+hwyMpg)/2) %>%  
  arrange(avgMpg) -> japaneseAutos
```

3. PIPING DATA

READING PIPES

- ▶ Pipes can be read in sequential order:

1. Take the mpg data frame, then
2. select the manufacturer, model, and fuel efficiency variables, then
3. rename the city gas mileage variable, then
4. rename the highway gas mileage variable, then
5. filter observations for Japanese automobile manufacturers, then
6. create a new average miles per gallon variable, then
7. arrange observations from high to low based on the new fuel efficiency variable, then
8. assign these changes to a new data frame named japaneseAutos

```
> library(tidyverse)
> mpg %>%
  select(manufacturer,
         model, cty, hwy) %>%
  rename(cityMpg = cty) %>%
  rename(hwyMpg = hwy) %>%
  filter(manufacturer == "honda" |
         manufacturer == "nissan" |
         manufacturer == "subaru" |
         manufacturer == "toyota") %>%
  mutate(avgMpg =
        (cityMpg+hwyMpg)/2) %>%
  arrange(avgMpg) -> japaneseAutos
```

3. PIPING DATA

READING PIPES

- ▶ Pipes can be read in sequential order:

1. Take the mpg data frame, **then**
2. select the manufacturer, model, and fuel efficiency variables, **then**
3. rename the city gas mileage variable, **then**
4. rename the highway gas mileage variable, **then**
5. filter observations for Japanese automobile manufacturers, **then**
6. create a new average miles per gallon variable, **then**
7. arrange observations from high to low based on the new fuel efficiency variable, **then**
8. assign these changes to a new data frame named japaneseAutos

```
> library(tidyverse)
> mpg %>%
  select(manufacturer,
         model, cty, hwy) %>%
  rename(cityMpg = cty) %>%
  rename(hwyMpg = hwy) %>%
  filter(manufacturer == "honda" |
         manufacturer == "nissan" |
         manufacturer == "subaru" |
         manufacturer == "toyota") %>%
  mutate(avgMpg =
        (cityMpg+hwyMpg)/2) %>%
  arrange(avgMpg) -> japaneseAutos
```

READING PIPES

- ▶ The final assignment can also be made on the first line of code like the example to the right
- ▶ I prefer the initial method only because the code “reads” in a linear fashion
- ▶ In either case, the data reference in each function can be omitted since it is “passed” by the pipe operator
- ▶ Pipes should be *short*

```
> library(tidyverse)
> japaneseAutos <- mpg %>%
  select(manufacturer,
         model, cty, hwy) %>%
  rename(cityMpg = cty) %>%
  rename(hwyMpg = hwy) %>%
  filter(manufacturer == "honda" |
         manufacturer == "nissan" |
         manufacturer == "subaru" |
         manufacturer == "toyota") %>%
  mutate(avgMpg =
        (cityMpg+hwyMpg)/2) %>%
  arrange(avgMpg)
```

PIPES AND GGPLOT2

- ▶ If we remove the data assignment, pipes still work!
- ▶ They will temporarily alter the data without making those changes permanent
- ▶ This is perfect behavior for making ggplot plots on a modified set of data without creating a new data frame
- ▶ Note that the data reference is not needed in the ggplot function

```
> library(tidyverse)  
  
> mpg %>%  
  select(manufacturer,  
         model, cty, hwy) %>%  
  rename(cityMpg = cty) %>%  
  rename(hwyMpg = hwy) %>%  
  filter(manufacturer == "honda" |  
         manufacturer == "nissan" |  
         manufacturer == "subaru" |  
         manufacturer == "toyota") %>%  
  mutate(avgMpg =  
        (cityMpg+hwyMpg)/2) %>%  
  arrange(avgMpg) %>%  
  ggplot() +  
    geom_histogram(mapping =  
                  aes(avgMpg))
```